


Microsoft Azure

 databricks

Search data, notebooks, recents, and more...

CTRL + P

kumaran

New

Workspace

Recents

Catalog

Workflows

Compute

Data Engineering

Job Runs

Machine Learning

Playground

Experiments

Features

Models

Serving

Partner Connect

Untitled Notebook 2024-12-07 22:37:25

Python

File Edit View Run Help

Last edit was 23 minutes ago

Run all

PRANAV-007

Schedule

Share

Workspace

palanisriram316@gmail.com

reviews.csv

Untitled Notebook 2024-12-07 19:02:45

Untitled Notebook 2024-12-07 22:37:25

31 minutes ago (1s)

1

```
from azure.storage.blob import BlobServiceClient
import io
import pandas as pd

# Connection string to your Azure Blob Storage account
connection_string = "DefaultEndpointsProtocol=https;AccountName=pranav67;
AccountKey=Y20CeHmVP6neuPVkn/
Vp2Nskrn8DbfxWe4kIHjAGSwbDX3ljxcKPefh0eAnadpXZQaosxNAI8ZDz+AStqFtB7w==;
EndpointSuffix=core.windows.net"

# Define the container name and the blob (file) name
container_name = "velan"
blob_name = "dataset01.csv"

# Initialize the BlobServiceClient using the connection string
blob_service_client = BlobServiceClient.from_connection_string(
    connection_string)

# Get a BlobClient for the blob you want to read
blob_client = blob_service_client.get_blob_client(container=container_name,
    blob=blob_name)

# Download the blob's content into memory (as a stream)
blob_data = blob_client.download_blob()

# Read the content of the blob as a string (this will be CSV data)
csv_data = blob_data.readall()

# Convert the CSV data into a pandas DataFrame for easy manipulation
data_frame = pd.read_csv(io.BytesIO(csv_data))

# Print the DataFrame to see the contents
print(data_frame)
```

	product_id	product_name	customer_location	target_column
0	1011	Tablet	Miami	1
1	1012	Smartwatch	New York	1
2	1013	Microwave	Nevada	0
3	1014	Wireless Earbuds	Chicago	1
4	1015	Desk Chair	Los Angeles	0
5	1016	Air Conditioner	Phoenix	1
6	1017	Soundbar	Dallas	1
7	1018	Dryer	Seattle	0
8	1019	Toaster	San Francisco	1
9	1020	VR Headset	Austin	1

[10 rows x 9 columns]

Last execution failed

2

```
1 import pandas as pd
2 import numpy as np
3
4 # Assuming data_frame is your cleaned DataFrame
5
6 # 1. Cumulative Sum for 'quantity_sold'
7 data_frame['cumulative_quantity_sold'] = data_frame['quantity_sold'].cumsum()
8
9 # 2. Rolling Averages
10 # We can calculate rolling averages for numerical columns like 'price',
11 # 'quantity_sold', and 'discount'
12 # For simplicity, we use a 3-row rolling window for each of these columns
13 data_frame['rolling_avg_price'] = data_frame['price'].rolling(window=3).mean()
14 data_frame['rolling_avg_quantity_sold'] = data_frame['quantity_sold'].rolling(window=3).mean()
15 data_frame['rolling_avg_discount'] = data_frame['discount'].rolling(window=3).mean()
```


```


(window=3).mean()
16
17 # 3. Date-based Features (if you have a 'date' column)
18 # For this example, assume the date column exists and is named 'date'
19 # You can extract year, month, day, weekday, etc.
20 # If there's no date column, you can skip this step
21
22 # Example: Assuming 'date' is a datetime column
23 if 'date' in data_frame.columns:
24     data_frame['year'] = data_frame['date'].dt.year
25     data_frame['month'] = data_frame['date'].dt.month
26     data_frame['day'] = data_frame['date'].dt.day
27     data_frame['weekday'] = data_frame['date'].dt.weekday
28
29 # 4. One-Hot Encoding for 'category' and 'customer_location'
30 # One-Hot Encoding using pandas' get_dummies
31 data_frame = pd.get_dummies(data_frame, columns=['category',
32     'customer_location'], drop_first=True)
33
34 # 5. Creating Interaction Features
35 # For example, the interaction between 'price' and 'quantity_sold' might
36 # be useful
37 data_frame['price_quantity_interaction'] = data_frame['price'] * data_frame
38     ['quantity_sold']
39
40 # 6. Feature Transformation: Log Transformation (for highly skewed
41 # features)
42 # Apply log transformation to 'price' if it is highly skewed
43 if data_frame['price'].skew() > 1: # check if skew is significant
44     data_frame['log_price'] = np.log1p(data_frame['price']) # log1p
45     handles 0 values safely
46
47 # Apply log transformation to 'quantity_sold' if needed
48 if data_frame['quantity_sold'].skew() > 1:
49     data_frame['log_quantity_sold'] = np.log1p(data_frame['quantity_sold'])
50
51 # 7. Additional Custom Feature (e.g., price per unit sold)
52 data_frame['price_per_unit'] = data_frame['price'] / (data_frame
53     ['quantity_sold'] + 1)
54
55 # 8. Remove any unnecessary columns (optional)
56 # Example: If 'product_name' and 'target_column' aren't useful for
57 # modeling, drop them
58 data_frame = data_frame.drop(columns=['product_name', 'target_column'])
59
60 # Show the final engineered DataFrame
61 print(data_frame)
62 X = X.dropna()
63 y = y[X.index] # Ensure that y matches the rows in X
64
65 # Split the data into training and testing sets again after dropping NaN
66 # rows
67 from sklearn.model_selection import train_test_split
68 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
69     random_state=42)

```


	product_id	price	...	price_quantity_interaction	price_per_unit
0	1011	750	...	13500	39.473684
1	1012	250	...	12500	4.901961
2	1013	120	...	2400	5.714286
3	1014	180	...	5040	6.206897
4	1015	150	...	3300	6.521739
5	1016	500	...	7500	31.250000
6	1017	350	...	12250	9.722222
7	1018	800	...	8000	72.727273
8	1019	60	...	2400	1.463415
9	1020	320	...	8000	12.307692

[10 rows x 23 columns]

 > TypeError: 'NoneType' object is not subscriptable

 Diagnose error

 Debug

Assistant Quick Fix: ON 

►  Last execution failed 3

```

1 from sklearn.impute import SimpleImputer
2 from sklearn.model_selection import train_test_split, cross_val_score,
  KFold
3 from sklearn.linear_model import LinearRegression

```

```

4 from sklearn.metrics import mean_squared_error, mean_absolute_error,
  r2_score
5 import numpy as np
6
7 # Assuming 'X' is the feature matrix and 'y' is the target variable
8 # Impute missing values in X
9 imputer = SimpleImputer(strategy='mean')
10 X_imputed = imputer.fit_transform(X)
11
12 # Split the data into training and testing sets
13 X_train, X_test, y_train, y_test = train_test_split(X_imputed, y,
  test_size=0.2, random_state=42)
14
15 # Train the Linear Regression model
16 model = LinearRegression()
17 model.fit(X_train, y_train)
18
19 # Make predictions on the test set
20 y_pred = model.predict(X_test)
21
22 # Evaluate the model using RMSE and R² on the test set
23 mae = mean_absolute_error(y_test, y_pred)
24 print(f"Mean Absolute Error (MAE): {mae}")
25
26 mse = mean_squared_error(y_test, y_pred)
27 print(f"Mean Squared Error (MSE): {mse}")
28
29 rmse = np.sqrt(mse)
30 print(f"Root Mean Squared Error (RMSE): {rmse}")
31
32 r2 = r2_score(y_test, y_pred)
33 print(f"R-squared (R²): {r2}")
34
35 # Cross-validation evaluation using RMSE and R²
36 kf = KFold(n_splits=5, shuffle=True, random_state=42)
37
38 # Using built-in scoring methods for cross-validation
39 cv_rmse_scores = cross_val_score(model, X_imputed, y, cv=kf,
  scoring='neg_root_mean_squared_error')
40 cv_r2_scores = cross_val_score(model, X_imputed, y, cv=kf, scoring='r2')
41
42 # Convert negative RMSE to positive values
43 cv_rmse_scores = -cv_rmse_scores
44
45 print(f"Cross-Validation RMSE Scores: {cv_rmse_scores}")
46 print(f"Mean CV RMSE: {cv_rmse_scores.mean()}")
47
48 print(f"Cross-Validation R² Scores: {cv_r2_scores}")
49 print(f"Mean CV R²: {cv_r2_scores.mean()}")

```

```

❶ > ValueError: Cannot use mean strategy with non-numeric data:
could not convert string to float: 'Microwave'
File <command-2434239681457285>, line 10
    7 # Assuming 'X' is the feature matrix and 'y' is the target variable
    8 # Impute missing values in X
    9 imputer = SimpleImputer(strategy='mean')
--> 10 X_imputed = imputer.fit_transform(X)
    12 # Split the data into training and testing sets
    13 X_train, X_test, y_train, y_test = train_test_split(X_imputed, y, test_s
ize=0.2, random_state=42)
~
File /databricks/python/lib/python3.11/site-packages/sklearn/impute/_base.py:32
7, in SimpleImputer._validate_input(self, X, in_fit)
    321 if "could not convert" in str(ve):
    322     new_ve = ValueError(
    323         "Cannot use {} strategy with non-numeric data:\n{}".format(
    324             self.strategy, ve
    325         )
    326     )
--> 327 raise new_ve from None
    328 else:
    329     raise ve

```

🔗 Diagnose error

🔍 Debug

Assistant Quick Fix: ON ▼

▶ ✓ 23 minutes ago (1s)

4

```

import joblib
from azure.storage.blob import BlobServiceClient

```

```

import os
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# Assuming 'X' is the feature matrix and 'y' is the target variable
# Example: Train a linear regression model (replace with your model and
training logic)
# Here, I use random data as an example. Replace this with your actual
dataset.
X = np.random.rand(100, 5) # 100 samples, 5 features
y = np.random.rand(100) # 100 target values

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Train a Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Evaluate the model
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse}")

# Save the trained model to a local file
model_filename = "linear_regression_model.pkl"
joblib.dump(model, model_filename)

# Azure Blob Storage details
connection_string = "DefaultEndpointsProtocol=https;AccountName=pranav67;
AccountKey=Y20CeHmVP6neuPVkn/
Vp2Nskrn8DbfxWe4kIHjAGSwbDX3ljxcKPeFh0eAnadpXZQAosxNAI8ZDz+ASTqFtB7w==;
EndpointSuffix=core.windows.net"
container_name = "velan"
model_blob_name = "linear_regression_model.pkl" # Blob name in the container

# Initialize the BlobServiceClient using the connection string
blob_service_client = BlobServiceClient.from_connection_string
(connection_string)

# Get the BlobClient for the model file
blob_client = blob_service_client.get_blob_client(container=container_name,
blob=model_blob_name)

# Upload the model file to Azure Blob Storage
try:
    # Open the model file in binary mode and upload to Blob Storage
    with open(model_filename, "rb") as data:
        blob_client.upload_blob(data, overwrite=True) # overwrite=True to
replace any existing file

    print(f"Model successfully uploaded to Azure Blob Storage as
{model_blob_name}")

    # Optionally, remove the local model file after uploading
    os.remove(model_filename)

except Exception as e:
    print(f"Error uploading model to Azure Blob Storage: {e}")

```

```

Mean Squared Error (MSE): 0.09808504910703829
Model successfully uploaded to Azure Blob Storage as linear_regression_model.pkl

```



5

[Shift+Enter] to run and move to next cell  
[Ctrl+Shift+P] to open the command palette  
[Esc H] to see all keyboard shortcuts

